

Lab 1

李秋宇 3220103373

Design

4.1 准备工程

从仓库拉取了源码

4.2 RV64内核引导

4.2.1 Makefile

在这里编写 `printk.c` 的编译规则

设定 `clean` 规则以完成清除构建文件

注意到 `$(CFLAG)` 这些变量已经在最顶层的 `Makefile` 中定义，因此直接使用

Makefile	lib/Makefile
1	<code>SRC = \$(wildcard *.c)</code>
2	<code>OBJ = \$(patsubst %.c,%.o,\$(SRC))</code>
3	
4	<code>all: \$(OBJ)</code>
5	
6	<code>%.o: %.c</code>
7	<code>\$(GCC) \$(CFLAG) -c \$< -o \$@</code>
8	
9	<code>.PHONY: clean</code>
10	<code>clean:</code>
11	<code>rm -f *.o</code>

完成

4.2.2 head.S

根据要求写汇编

注意到这里已经定义了 `start_kernel` 并且有栈和栈顶的信息，直接使用即可

ARM Assembly	head.S
1	<code>.extern start_kernel</code>

```

2      .section .text.entry
3      .globl _start
4      _start:
5          la sp, boot_stack_top # set stack pointer to top of boot stack
6          jal start_kernel # jump to start_kernel
7
8      .section .bss.stack
9      .globl boot_stack
10     boot_stack:
11         .space 4096 # Stack size 4KiB
12
13     .globl boot_stack_top
14     boot_stack_top:

```

4.2.3 sbi.c

参照示例进行内联汇编撰写，照葫芦画瓢

使用伪指令 `mv` 把寄存器拷贝，然后按照内联汇编格式写，写完返回结构体就完事儿了兄弟们

C *sbi.c*

```

1  struct sbiret sbi_ecall(uint64_t eid, uint64_t fid,
2                          uint64_t arg0, uint64_t arg1, uint64_t arg2,
3                          uint64_t arg3, uint64_t arg4, uint64_t arg5) {
4      struct sbiret ret;
5      asm volatile (
6          "mv a7, %[eid]\n"
7          "mv a6, %[fid]\n"
8          "mv a0, %[arg0]\n"
9          "mv a1, %[arg1]\n"
10         "mv a2, %[arg2]\n"
11         "mv a3, %[arg3]\n"
12         "mv a4, %[arg4]\n"
13         "mv a5, %[arg5]\n"
14         "ecall\n"
15         "mv %[error], a0\n"
16         "mv %[value], a1\n"
17         : [error] "=r" (ret.error), [value] "=r" (ret.value)
18         : [eid] "r" (eid), [fid] "r" (fid),
19           [arg0] "r" (arg0), [arg1] "r" (arg1), [arg2] "r" (arg2),
20           [arg3] "r" (arg3), [arg4] "r" (arg4), [arg5] "r" (arg5)
21     );
22     return ret;
23 }

```

完成 `sbi_ecall` 后只需要在剩下几个函数里面调用就完事儿了兄弟们

根据表格中每个函数对应的 `eid` 和 `fid` 直接调用

C *sbi.c*

```

1  struct sbiret sbi_set_timer(uint64_t stime_value) {
2      struct sbiret ret;
3      ret = sbi_ecall(0x54494d45, 0x0, stime_value, 0, 0, 0, 0, 0);
4      return ret;
5  }
6
7  struct sbiret sbi_debug_console_write_byte(uint8_t byte) {
8      struct sbiret ret;
9      ret = sbi_ecall(0x4442434E, 0x2, byte, 0, 0, 0, 0, 0);
10     return ret;
11 }
12
13 struct sbiret sbi_system_reset(uint32_t reset_type, uint32_t reset_reason)
14 {
15     struct sbiret ret;
16     ret = sbi_ecall(0x53525354, 0x0, reset_type, reset_reason, 0, 0, 0,
17                     0);
18     return ret;
19 }

```

4.2.4 defs.h

依旧是照葫芦画瓢，使用指令 `csrr` 进行读取

C *defs.h*

```

1  #define csr_read(csr) \
2      ({ \
3          uint64_t __v; \
4          asm volatile("csrr %0, " #csr : "=r"(__v)); \
5          __v; \
6      })

```

Run

完成上述后，进行构建

```

rm -f *.o
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/lib'
make -C init clean
make[1]: Entering directory '/home/lqy/ZJU/lab1/src/init'
make[1]: 'clean' is up to date.
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/init'
make -C arch/riscv clean
make[1]: Entering directory '/home/lqy/ZJU/lab1/src/arch/riscv'
make -C kernel clean
make[2]: Entering directory '/home/lqy/ZJU/lab1/src/arch/riscv/kernel'
make[2]: 'clean' is up to date.
make[2]: Leaving directory '/home/lqy/ZJU/lab1/src/arch/riscv/kernel'
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/arch/riscv'

Clean Finished
make -C lib all
make[1]: Entering directory '/home/lqy/ZJU/lab1/src/lib'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/lqy/ZJU/lab1/src/include -I /home/lqy/ZJU/lab1/src/arch/riscv/include -c printk.c -o printk.o
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/lib'
make -C init all
make[1]: Entering directory '/home/lqy/ZJU/lab1/src/init'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/lqy/ZJU/lab1/src/include -I /home/lqy/ZJU/lab1/src/arch/riscv/include -c main.c
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/lqy/ZJU/lab1/src/include -I /home/lqy/ZJU/lab1/src/arch/riscv/include -c test.c
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/init'
make -C arch/riscv all
make[1]: Entering directory '/home/lqy/ZJU/lab1/src/arch/riscv'
make -C kernel all
make[2]: Entering directory '/home/lqy/ZJU/lab1/src/arch/riscv/kernel'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/lqy/ZJU/lab1/src/include -I /home/lqy/ZJU/lab1/src/arch/riscv/include -c head.S
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/lqy/ZJU/lab1/src/include -I /home/lqy/ZJU/lab1/src/arch/riscv/include -c sbi.c
make[2]: Leaving directory '/home/lqy/ZJU/lab1/src/arch/riscv/kernel'
riscv64-linux-gnu-ld -T kernel/vmlinux.lds kernel/*.o .././init/*.o .././lib/*.o -o .././vmlinux
riscv64-linux-gnu-objcopy -O binary .././vmlinux ./boot/image
riscv64-linux-gnu-objdump -S .././vmlinux > .././vmlinux.asm
nm .././vmlinux > .././System.map
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/arch/riscv'

Build Finished OK
(base) lqy@LAPTOP-SPV26TG8:~/ZJU/src

```

构建完成，使用 `make run` 进行运行

```

Platform IPI Device      : aclint-mswi
Platform Timer Device    : aclint-mtimer @ 100000000Hz
Platform Console Device  : uart8250
Platform HSM Device      : ---
Platform PMU Device      : ---
Platform Reboot Device   : sifive_test
Platform Shutdown Device : sifive_test
Platform Suspend Device  : ---
Platform CPPC Device     : ---
Firmware Base           : 0x80000000
Firmware Size           : 322 KB
Firmware RW Offset      : 0x40000
Firmware RW Size        : 66 KB
Firmware Heap Offset    : 0x48000
Firmware Heap Size      : 34 KB (total), 2 KB (reserved), 9 KB (used), 22 KB (free)
Firmware Scratch Size   : 4096 B (total), 760 B (used), 3336 B (free)
Runtime SBI Version     : 1.0

Domain0 Name            : root
Domain0 Boot HART       : 0
Domain0 HARTs           : 0*
Domain0 Region00        : 0x0000000002000000-0x000000000200ffff M: (I,R,W) S/U: ()
Domain0 Region01        : 0x00000000080040000-0x0000000008005ffff M: (R,W) S/U: ()
Domain0 Region02        : 0x00000000080000000-0x0000000008003ffff M: (R,X) S/U: ()
Domain0 Region03        : 0x00000000000000000-0xffffffffffffffff M: (R,W,X) S/U: (R,W,X)
Domain0 Next Address    : 0x00000000080200000
Domain0 Next Arg1       : 0x00000000087e00000
Domain0 Next Mode       : S-mode
Domain0 SysReset        : yes
Domain0 SysSuspend      : yes

Boot HART ID            : 0
Boot HART Domain        : root
Boot HART Priv Version  : v1.12
Boot HART Base ISA      : rv64imafdch
Boot HART ISA Extensions : time,sstc
Boot HART PMP Count     : 16
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count    : 16
Boot HART MIDELEG       : 0x00000000000001666
Boot HART MEDELEG       : 0x00000000000f0b509
2024 ZJU Operating System
(base) lqv@LAPTOP-SDV26TG8:~/Z/1/src>

```

完成!

4.3 RV64时钟中断处理

这一部分写其实压栈了，因为相当于先写4.3.3再写4.3.2再倒着写 🤔

4.3.0 Preparation

首先根据准备工作完成对应修改

4.3.1 head.S

除了根据模板作一些调整之外，还编写了以下四部曲

1. 根据 `_traps` 地址赋值给 `stvec` 用于中断处理

2. 设置开启时钟中断，把 `sie` 的第二位置 1
3. 设置第一个时钟中断，这里根据了 `clock.c` 里面的定义设定了下一次时钟中断
4. 开启中断响应

ARM Assembly *head.S*

```

1  __start:
2
3      # set stvec = _traps
4      la t0, _traps
5      csrw stvec, t0
6
7      # set sie[STIE] = 1
8      csrr t0, sie
9      li t1, 0x20
10     or t0, t0, t1
11     csrw sie, t0
12
13     # set first time interrupt
14     li t0, 0x10000000
15     csrw stimecmp, t0
16
17     # set sstatus[SIE] = 1
18     csrr t0, sstatus
19     li t1, 0x2
20     or t0, t0, t1
21     csrw sstatus, t0

```

4.3.2 entry.S

还是四部曲

1. 首先压栈保存所有寄存器，注意到这里其实我没有保存 `x0`，因为感觉没啥必要🤔
2. 调用 `trap_handler`，这里后来注意到还需要传参，通过 `a0` 和 `a1` 进行传参
3. 上下文恢复，注意 `sp` 最后恢复，顺序和保存时相反
4. 返回

ARM Assembly *entry.S*

```

1      .extern trap_handler
2      .section .text.entry
3      .align 2
4      .globl _traps
5  _traps:
6      # 1. save 32 registers and sepc to stack
7      addi sp, sp, -32*8
8      sd ra, 0(sp) # x1
9      sd sp, 8(sp) # x2

```

```

10     sd gp, 16(sp) # x3
11     sd tp, 24(sp) # x4
12     sd t0, 32(sp) # x5
13     sd t1, 40(sp) # x6
14     sd t2, 48(sp) # x7
15     sd s0, 56(sp) # x8
16     sd s1, 64(sp) # x9
17     sd a0, 72(sp) # x10
18     sd a1, 80(sp) # x11
19     sd a2, 88(sp) # x12
20     sd a3, 96(sp) # x13
21     sd a4, 104(sp) # x14
22     sd a5, 112(sp) # x15
23     sd a6, 120(sp) # x16
24     sd a7, 128(sp) # x17
25     sd s2, 136(sp) # x18
26     sd s3, 144(sp) # x19
27     sd s4, 152(sp) # x20
28     sd s5, 160(sp) # x21
29     sd s6, 168(sp) # x22
30     sd s7, 176(sp) # x23
31     sd s8, 184(sp) # x24
32     sd s9, 192(sp) # x25
33     sd s10, 200(sp) # x26
34     sd s11, 208(sp) # x27
35     sd t3, 216(sp) # x28
36     sd t4, 224(sp) # x29
37     sd t5, 232(sp) # x30
38     sd t6, 240(sp) # x31
39     csrr t0, sepc
40     sd t0, 248(sp) # sepc
41
42     # 2. call trap_handler
43     csrr a0, scause # arg1: scause
44     csrr a1, sepc # arg2: sepc
45     call trap_handler
46
47     # 3. restore sepc and 32 registers (x2(sp) should be restore last)
    from stack
48     ld t0, 248(sp) # sepc
49     csrwr sepc, t0
50     ld t6, 240(sp) # x31
51     ld t5, 232(sp) # x30
52     ld t4, 224(sp) # x29
53     ld t3, 216(sp) # x28
54     ld s11, 208(sp) # x27
55     ld s10, 200(sp) # x26
56     ld s9, 192(sp) # x25
57     ld s8, 184(sp) # x24
58     ld s7, 176(sp) # x23

```

```

59     ld s6, 168(sp) # x22
60     ld s5, 160(sp) # x21
61     ld s4, 152(sp) # x20
62     ld s3, 144(sp) # x19
63     ld s2, 136(sp) # x18
64     ld a7, 128(sp) # x17
65     ld a6, 120(sp) # x16
66     ld a5, 112(sp) # x15
67     ld a4, 104(sp) # x14
68     ld a3, 96(sp)  # x13
69     ld a2, 88(sp)  # x12
70     ld a1, 80(sp)  # x11
71     ld a0, 72(sp)  # x10
72     ld s1, 64(sp)  # x9
73     ld s0, 56(sp)  # x8
74     ld t2, 48(sp)  # x7
75     ld t1, 40(sp)  # x6
76     ld t0, 32(sp)  # x5
77     ld tp, 24(sp)  # x4
78     ld gp, 16(sp)  # x3
79     ld ra, 0(sp)   # x1
80     ld sp, 8(sp)   # x2
81     addi sp, sp, 32*8
82
83     # 4. return from trap
84     sret

```

4.3.3 trap.c

查阅了[priv-isa-asciidoc第43页](#)可以得到对应的异常代码，然后判断 `scause` 是否符合即可

注意到还要输出，所以调用 `printk` 配合输出

4.3.4 clock.c

主要是内联汇编，根据提示写

顺便调用一下前面已经实现的 `sbi_set_timer`

C *clock.c*

```

1  #include "../include/stdint.h"
2  #include "sbi.h"
3
4  // QEMU 中时钟的频率是 10MHz，也就是 1 秒钟相当于 100000000 个时钟周期
5  uint64_t TIMECLOCK = 100000000;
6
7  uint64_t get_cycles() {

```



```
8      // 编写内联汇编，使用 rdttime 获取 time 寄存器中（也就是 mtime 寄存器）的值并
      返回
9      uint64_t time;
10     asm volatile ("rdtime %0" : "=r" (time));
11     return time;
12 }
13
14 void clock_set_next_event() {
15     // 下一次时钟中断的时间点
16     uint64_t next = get_cycles() + TIMECLOCK;
17
18     // 使用 sbi_set_timer 来完成对下一次时钟中断的设置
19     sbi_set_timer(next);
20     return ;
21 }
```

4.3.5 test.c

根据文档修改

编译成功

运行成功

[CS] Supervisor mode timer interrupt!

1

请总结一下RISC-V的calling convention, 并解释Caller / Callee Saved Register有什么区别?

Calling Convention:

- 函数传参使用寄存器 **a0** 到 **a7** , 如果超过8个则使用栈
- 函数返回值使用寄存器 **a0** 和 **a1** 返回
- 函数调用时通常需要保存当前的上下文并调整栈指针

Caller Saved Register: 调用者保存寄存器, 在RISC-V中为 **t0** 到 **t6** , 在调用函数之前需要手动保存

Callee Saved Register: 被调用者保存寄存器, 在RISC-V中为 **s0** 到 **s11** , 在被调用者进入函数时保存在栈上, 在函数返回前恢复

2

编译之后, 通过System.map查看vmlinux.lds中自定义符号的值并截图。

```
riscv64-linux-gnu-ld -T kernel/vmlinux.lds kernel/*.o ../../init/*.o ../../lib/*.o -o ../../vmlinux
riscv64-linux-gnu-objcopy -O binary ../../vmlinux ./boot/Image
riscv64-linux-gnu-objdump -S ../../vmlinux > ../../vmlinux.asm
rm ../../vmlinux > ../../System.map
make[1]: Leaving directory '/home/lqy/ZJU/lab1/src/arch/riscv'

Build Finished OK
(base) lqy@LAPTOP-SDV26TG8:~/Z/l/src$ ls
Makefile  System.map  arch/  fw_jump.bin  include/  init/  lib/  vmlinux*  vmlinux.asm
(base) lqy@LAPTOP-SDV26TG8:~/Z/l/src$ rm vmlinux
0000000000200000 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200048 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200168 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200190 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
00000000002001d8 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200294 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200330 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
00000000002003d0 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
000000000020047c t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200504 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200548 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200598 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
00000000002005e0 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200640 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
00000000002008ac t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200934 t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200c3c t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
000000000020142c t $xrv64i2p1_m2p0_a2p1_f2p2_d2p2_zicsr2p0_zmmulip0
0000000000200000 A BASE_ADDR
0000000000203000 D TIMECLOCK
0000000000203008 d _GLOBAL_OFFSET_TABLE_
0000000000205000 B _ebss
0000000000203008 D _edata
0000000000205000 B _kernel
00000000002020e9 R _erodata
00000000002014ac T _etext
0000000000204000 B _sbss
0000000000203000 D _sdata
0000000000200000 T _kernel
0000000000203000 R _srodata
```

3

用 `csr_read` 宏读取 `sstatus` 寄存器的值, 对照RISC-V手册解释其含义并截图

在 `arch/riscv/kernel/trap.c` 中加入读取 `sstatus` 的逻辑:

```
C trap.c

1 uint64_t status = csr_read(sstatus);
2 printk("sstatus = %lx\n", status);
```

重新编译运行后得到

```

kernel is running!
kernel is running!
kernel is running!
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
kernel is running!
QEMU: Terminated

```

即 `sstatus = 0x8000000200006120`，即二进制

Verilog

```
1 64'b10000000_00000000_00000000_00000010_00000000_00000000_01100001_00100000
```

对应置 1 位的具体信息为

- SPIE (5): 上一次进入S模式之前的中断使能状态
- SPP (8): S模式中断处理前的特权模式为机器模式
- FS (13,14)
- UXL[1] (33)
- SD (63)

4

用 `csr_write` 宏向 `sscratch` 寄存器写入数据，并验证是否写入成功并截图

还是在 `trap.c` 中写

C `trap.c`

```

1 csr_write(sscratch, 0x666);
2 uint64_t res = csr_read(sscratch);
3 printf("sscratch = %lx\n", res);

```

编译后运行结果为

```
kernel is running!
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
kernel is running!
[S] Supervisor mode timer interrupt!
sstatus = 8000000200006120
sscratch = 666
QEMU: Terminated
```

5

详细描述你可以通过什么步骤来得到 `arch/arm64/kernel/sys.i`，给出过程以及截图

首先安装交叉编译链工具 `gcc-aarch64-linux-gnu`

然后使用命令：

Shell

```
1 aarch64-linux-gnu-gcc -E init/main.c -o arch/arm64/kernel/sys.i -Iinclude
```

即可生成

```

(base) lqy@LAPTOP-SDV26TG8:~/Z/l/src>aarch64-linux-gnu-gcc -E init/main.c -o arch/arm64/kernel/sys.i
init/main.c:1:10: fatal error: printk.h: No such file or directory
   1 | #include "printk.h"
     |           ^~~~~~
compilation terminated.
(base) lqy@LAPTOP-SDV26TG8:~/Z/l/src[1]>aarch64-linux-gnu-gcc -E init/main.c -o arch/arm64/kernel/sys.i -Iinclude
(base) lqy@LAPTOP-SDV26TG8:~/Z/l/src>

```

C *arch/arm64/kernel/sys.i*

```

1  # 0 "init/main.c"
2  # 0 "<built-in>"
3  # 0 "<command-line>"
4  # 1 "/usr/aarch64-linux-gnu/include/stdc-predef.h" 1 3
5  # 0 "<command-line>" 2
6  # 1 "init/main.c"
7  # 1 "include/printk.h" 1
8
9
10
11 # 1 "include/stddef.h" 1
12
13
14
15 typedef long int ptrdiff_t;
16 typedef long unsigned int size_t;
17 typedef unsigned int wchar_t;
18
19
20
21
22 typedef __builtin_va_list va_list;
23 # 5 "include/printk.h" 2
24
25
26
27
28
29 int printk(const char *, ...);
30 # 2 "init/main.c" 2
31
32 extern void test();
33
34 int start_kernel() {
35     printk("2024");
36     printk(" ZJU Operating System\n");
37
38     test();
39     return 0;
40 }

```

6

寻找 Linux v6.0 中 ARM32 RV32 RV64 x86_64 架构的系统调用表；请列出源代码文件，展示完整的系统调用表（宏展开后），每一步都需要截图。

🔗 方便

由于本人之前使用了Linux 6.11，因此这里为了方便没有使用Linux v6.0而是直接在6.11基础上进行这一步

各个架构的系统调用表都在Linux源码的根目录下的 `arch/` 中

Shell ARM32

```
1 $ cd arch/arm/include/asm
2 $ ls | grep syscall
3 syscall.h
4 syscalls.h
```

```
• (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/a/i/asm➤pwd
/home/lqy/ZJU/Rinux/linux-6.11/arch/arm/include/asm
• (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/a/i/asm➤ls |grep syscall
syscall.h
syscalls.h
○ (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/a/i/asm➤
```

Shell RV

```
1 $ cd arch/riscv/kernel
2 $ ls | grep syscall
3 Makefile.syscalls
4 compat_syscall_table.c
5 compat_syscall_table.o
6 syscall_table.c
7 syscall_table.o
```

```
• (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/r/kernel➤ls |grep syscall
Makefile.syscalls
compat_syscall_table.c
compat_syscall_table.o
syscall_table.c
syscall_table.o
○ (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/r/kernel➤
```

Shell x86_64

```
1 $ cd arch/x86/entry/syscalls
2 $ ls
3 Makefile syscall_32.tbl syscall_64.tbl
```

```
• (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/x/e/syscalls>ls
Makefile syscall_32.tbl syscall_64.tbl
○ (base) lqy@LAPTOP-SDV26TG8:~/Z/R/l/a/x/e/syscalls>
```

```
lab1.md fish syscalls syscall_64.tbl X
home > lqy > ZJU > Rlinux > linux-6.11 > arch > x86 > entry > syscalls > syscall_64.tbl
1 # SPDX-License-Identifier: GPL-2.0 WITH Linux-syscall-note
2 #
3 # 64-bit system call numbers and entry vectors
4 #
5 # The format is:
6 # <number> <abi> <name> <entry point> [<compat entry point> [noreturn]]
7 #
8 # The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls
9 #
10 # The abi is "common", "64" or "x32" for this file.
11 #
12 0 common read sys_read
13 1 common write sys_write
14 2 common open sys_open
15 3 common close sys_close
16 4 common stat sys_newstat
17 5 common fstat sys_newfstat
18 6 common lstat sys_newlstat
19 7 common poll sys_poll
20 8 common lseek sys_lseek
21 9 common mmap sys_mmap
22 10 common mprotect sys_mprotect
23 11 common munmap sys_munmap
24 12 common brk sys_brk
25 13 64 rt_sigaction sys_rt_sigaction
26 14 common rt_sigprocmask sys_rt_sigprocmask
27 15 64 rt_sigreturn sys_rt_sigreturn
28 16 64 ioctl sys_ioctl
29 17 common pread64 sys_pread64
30 18 common pwrite64 sys_pwrite64
31 19 64 readv sys_readv
32 20 64 writev sys_writev
```

7

阐述什么是 ELF 文件？尝试使用 `readelf` 和 `objdump` 来查看 ELF 文件，并给出解释和截图。

ELF文件是一种广泛用于Unix和类Unix系统的可执行文件、目标代码、共享库的文件格式，其格式具有灵活性和可扩展性

- `readelf` 可以用于展示ELF文件的信息，有多种可选选项，包括 `-h` , `-a` 等
- `objdump` 可以用于展示目标文件的信息

Shell

```
1 $ readelf -h vmlinux
2 ELF Header:
3 Magic: 7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
4 Class: ELF64
```



```

5      Data:                2's complement, little endian
6      Version:             1 (current)
7      OS/ABI:              UNIX - System V
8      ABI Version:        0
9      Type:                EXEC (Executable file)
10     Machine:             RISC-V
11     Version:             0x1
12     Entry point address: 0x80200000
13     Start of program headers: 64 (bytes into file)
14     Start of section headers: 30864 (bytes into file)
15     Flags:               0x0
16     Size of this header:  64 (bytes)
17     Size of program headers: 56 (bytes)
18     Number of program headers: 4
19     Size of section headers: 64 (bytes)
20     Number of section headers: 20
21     Section header string table index: 19
22
23     $ objdump vmlinux -h
24
25     vmlinux:      file format elf64-little
26
27     Sections:
28     Idx Name          Size      VMA              LMA              File off
29     Algn
30     0 .text           000014fc  0000000080200000  0000000080200000  00001000
31     2**12
32     CONTENTS, ALLOC, LOAD, READONLY, CODE
33     1 .rodata         00000109  0000000080202000  0000000080202000  00003000
34     2**12
35     CONTENTS, ALLOC, LOAD, READONLY, DATA
36     2 .eh_frame       000002a8  0000000080202110  0000000080202110  00003110
37     2**3
38     CONTENTS, ALLOC, LOAD, READONLY, DATA
39     3 .data           00000008  0000000080203000  0000000080203000  00004000
40     2**12
41     CONTENTS, ALLOC, LOAD, DATA
42     4 .got            00000020  0000000080203008  0000000080203008  00004008
43     2**3
44     CONTENTS, ALLOC, LOAD, DATA
45     5 .got.plt        00000010  0000000080203028  0000000080203028  00004028
46     2**3
47     CONTENTS, ALLOC, LOAD, DATA
48     6 .bss            00001000  0000000080204000  0000000080204000  00004038
49     2**12
50     ALLOC
51     7 .debug_info     00000d12  0000000000000000  0000000000000000  00004038
52     2**0
53     CONTENTS, READONLY, DEBUGGING, OCTETS

```

```

45      8 .debug_abbrev 000005a5 0000000000000000 0000000000000000 00004d4a
      2**0
46          CONTENTS, READONLY, DEBUGGING, OCTETS
47      9 .debug_aranges 00000220 0000000000000000 0000000000000000 000052f0
      2**4
48          CONTENTS, READONLY, DEBUGGING, OCTETS
49     10 .debug_rnglists 0000010b 0000000000000000 0000000000000000 00005510
      2**0
50          CONTENTS, READONLY, DEBUGGING, OCTETS
51     11 .debug_line 00001379 0000000000000000 0000000000000000 0000561b
      2**0
52          CONTENTS, READONLY, DEBUGGING, OCTETS
53     12 .debug_str 0000035a 0000000000000000 0000000000000000 00006994
      2**0
54          CONTENTS, READONLY, DEBUGGING, OCTETS
55     13 .debug_line_str 00000116 0000000000000000 0000000000000000 00006cee
      2**0
56          CONTENTS, READONLY, DEBUGGING, OCTETS
57     14 .comment 00000026 0000000000000000 0000000000000000 00006e04
      2**0
58          CONTENTS, READONLY
59     15 .riscv.attributes 00000046 0000000000000000 0000000000000000
      00006e2a 2**0
60          CONTENTS, READONLY
61

```

运行一个 ELF 文件，然后通过 `cat /proc/PID/maps` 来给出其内存布局并截图。

这里随便写一个简单的C程序然后看

```

C      t.c

1      #include <stdio.h>
2
3      int main() {
4          int i = 0;
5          while (i < 1000000000000000000) {
6              printf("%d\n", i);
7              i++;
8          }
9          return 0;
10     }

```

编译并运行

Shell

```

1      gcc t.c -o a.out
2      ./a.out

```

这时开始疯狂跑循环输出

开另一个终端，查看PID

Shell

```

1  $ ps -e | grep a.out
2    90706 pts/5    00:00:04 a.out
3
4  $ cat /proc/90706/maps
5  5627216a8000-5627216a9000 r--p 00000000 08:20 303404
   /home/lqy/ZJU/lab1/a.out
6  5627216a9000-5627216aa000 r-xp 00001000 08:20 303404
   /home/lqy/ZJU/lab1/a.out
7  5627216aa000-5627216ab000 r--p 00002000 08:20 303404
   /home/lqy/ZJU/lab1/a.out
8  5627216ab000-5627216ac000 r--p 00002000 08:20 303404
   /home/lqy/ZJU/lab1/a.out
9  5627216ac000-5627216ad000 rw-p 00003000 08:20 303404
   /home/lqy/ZJU/lab1/a.out
10 562723442000-562723463000 rw-p 00000000 00:00 0
    [heap]
11 7f5712e0f000-7f5712e12000 rw-p 00000000 00:00 0
12 7f5712e12000-7f5712e3a000 r--p 00000000 08:20 291046
   /usr/lib/x86_64-linux-gnu/libc.so.6
13 7f5712e3a000-7f5712fc2000 r-xp 00028000 08:20 291046
   /usr/lib/x86_64-linux-gnu/libc.so.6
14 7f5712fc2000-7f5713011000 r--p 001b0000 08:20 291046
   /usr/lib/x86_64-linux-gnu/libc.so.6
15 7f5713011000-7f5713015000 r--p 001fe000 08:20 291046
   /usr/lib/x86_64-linux-gnu/libc.so.6
16 7f5713015000-7f5713017000 rw-p 00202000 08:20 291046
   /usr/lib/x86_64-linux-gnu/libc.so.6
17 7f5713017000-7f5713024000 rw-p 00000000 00:00 0
18 7f5713032000-7f5713034000 rw-p 00000000 00:00 0
19 7f5713034000-7f5713035000 r--p 00000000 08:20 291043
   /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
20 7f5713035000-7f5713060000 r-xp 00001000 08:20 291043
   /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
21 7f5713060000-7f571306a000 r--p 0002c000 08:20 291043
   /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
22 7f571306a000-7f571306c000 r--p 00036000 08:20 291043
   /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
23 7f571306c000-7f571306e000 rw-p 00038000 08:20 291043
   /usr/lib/x86_64-linux-gnu/ld-linux-x86-64.so.2
24 7ffe72256000-7ffe72277000 rw-p 00000000 00:00 0
    [stack]
25 7ffe722cf000-7ffe722d3000 r--p 00000000 00:00 0
    [vvar]

```

8

在我们使用 `make run` 时，OpenSBI 会产生如下输出：

```
Shell
```

```
1      OpenSBI v1.5.1
2
3      / _ \                               / ____| | _ \|   \_/_|
4    | | | | | _ \ _ _ _ _ _ _ _ _ _ _ | | (_____| |_| |
5    | | | | | ___ \___\___\___\___\___\_||_\__<|_| |
6    | | __| | |_) | _ _ _ _ _ _ _ _ _ _| | |_) | |_| |_
7    \___ _/ | . _ _/ \___ _|| | | _ _ _/ | _ _/ _ _ _|
8          ||
9         | _|
10
11     .....
12
13     Boot HART MIDELEG                   : 0x0000000000000222
14     Boot HART MEDELEG                    : 0x0000000000000b109
15
16     .....
```

通过查看[RISC-V Privileged Spec](#) 中的 `medeleg` 和 `mideleg` 部分, 解释上面 `MIDELEG` 和 `MEDELEG` 值的含义

在手册的**35**页有详细定义

通常来说，所有的**trap**都可以在**M**模式下解决，因为**M**模式下的陷入解决程序可以重定向到合适的权限等级进行解决陷入；但是为了提升性能表现，引入这两个寄存器

- **MIDELEG:** 用于保存哪些中断需要交给S模式处理
- **MEDELEG:** 用于保存哪些异常需要交给S模式处理

这里的 `MIDELEG == 0x222 == 0b001000100010`，即第1,5,9位置1，则表明

- Supervisor Software Interrupt
- Supervisor Timer Interrupt
- Supervisor External Interrupt

上述三种中断交给S模式处理

这里 `MEDELEG == 0xb109 == 0b101100001001`，即第0,3,8,9,11位置1，则表明

- Instruction Address Misaligned
- Breakpoint
- Environment Call from U-mode
- Environment Call from S-mode
- Environment Call from M-mode

上述五种异常交给S模式处理

Thinkings

这个lab总体来说难度不大，主要是学习一些新版本的知識，比如内联汇编之类的，而实验指导非常完善，夸夸助教们👍，因此上手很快，也很容易写完

内容不算多，而且学习新东西的过程比较有趣，所有总体顺利~

不过思考题觉得有点多，甚至觉得写思考题的时间比写lab还要就，所以偷懒了一些😁