

Lab 5

李秋宇 3220103373

Design

准备工程

添加文件，修改Makefile

缺页异常处理

实现虚拟内存管理功能

首先在 `proc.h` 中添加VMA的宏和数据结构

然后实现对 `struct vm_area_struct` 的查找和添加方法

注意到所有的VMA都是以链表的形式串在一起，所以

- 查找：从链表头部开始往后遍历，如果找到地址在某个VMA中就直接返回，否则一直往后遍历
- 添加：新建VMA，填入参数，插入到链表尾部

C `arch/riscv/kernel/proc.c`

```

1  struct vm_area_struct *find_vma(struct mm_struct *mm, uint64_t addr) {
2      struct vm_area_struct *vma = mm->mmap; // Header of the VMA list.
3      while (vma) {
4          if (addr ≥ vma->vm_start && addr < vma->vm_end) return vma;
5          vma = vma->vm_next;
6      }
7      return NULL;
8  }
9
10 uint64_t do_mmap(struct mm_struct *mm, uint64_t addr, uint64_t len, uint64_t
    vm_pgoff, uint64_t vm_filesz, uint64_t flags) {
11     struct vm_area_struct *vma = (struct vm_area_struct *)alloc_page();
12     vma->vm_start = PGROUNDDOWN(addr);
13     vma->vm_end = PGROUNDUP(addr + len);
14     vma->vm_pgoff = vm_pgoff;
15     vma->vm_filesz = vm_filesz;
16     vma->vm_flags = flags;
17
18     struct vm_area_struct *vma = (struct vm_area_struct *)alloc_page();
19     vma->vm_start = addr;
20     vma->vm_end = addr + len;
21     vma->vm_pgoff = vm_pgoff;
22     vma->vm_filesz = vm_filesz;
23     vma->vm_flags = flags;

```

```

24     vma->vm_next = NULL;
25
26     struct vm_area_struct *head = mm->mmap;
27     if (!head) {
28         mm->mmap = vma;
29         mm->mmap->vm_prev = NULL;
30     } else {
31         for (; head->vm_next; head = head->vm_next);
32         head->vm_next = vma;
33         vma->vm_prev = head;
34     }
35
36     return _sva;
37 }

```

修改 task_init

首先把原来的对用户程序uapp、用户态栈的空间分配和页表映射取消，改为Demand Paging模式

- 对于用户态栈来说，在VMA中映射区域为 [USER_END-PGSIZE:USER_END]，权限为R|W且为匿名区域
- 对于用户程序来说，在VMA中映射预取位 [vaddr:vaddr+memsz]，即它在内存中实际占据的空间，权限位R|W|X

C arch/riscv/kernel/proc.c

```

1  void load_program(struct task_struct *task) {
2      Elf64_Ehdr *ehdr = (Elf64_Ehdr *)_sramdisk;
3      Elf64_Phdr *phdrs = (Elf64_Phdr *)(_sramdisk + ehdr->e_phoff);
4      for (int i = 0; i < ehdr->e_phnum; ++i) {
5          Elf64_Phdr *phdr = phdrs + i;
6          if (phdr->p_type == PT_LOAD) {
7              do_mmap(&task->mm, phdr->p_vaddr, phdr->p_memsz, phdr->p_offset,
8                  phdr->p_filesz, VM_READ | VM_WRITE | VM_EXEC);
9          }
10     }
11     task->thread.sepc = ehdr->e_entry;
12 }
13 extern uint64_t swapper_pg_dir[];
14 void task_init() {
15     ...
16     for (int i = 1; i < NR_TASKS; i++) {
17         ...
18
19         // Create page table for user process.
20         ...
21
22         // Create VMA for user process.
23         task[i]->mm.mmap = NULL;
24
25         // Copy user app.

```

```

26         load_program(task[i]);
27
28         // Create stack for user process.
29         do_mmap(&task[i]→mm, USER_END - PGSIZE, PGSIZE, 0x0, 0x0, VM_READ |
VM_WRITE | VM_ANON);
30     }
31
32     printk(" ... task_init done!\n");
33 }

```

实现后截获的缺页异常符合实验指导中所述

实现 Page Fault Handler

在 `trap_handler` 中扩展缺页异常处理类型，都调用 `void do_page_fault(struct pt_regs *regs)`

C `arch/riscv/kernel/trap.c`

```

1  #define SUPERVISOR_TIMER_INTERRUPT_TYPE    0x8000000000000005
2  #define ENVIRONMRNT_CALL_FROM_U_MODE_TYPE 0x8
3  #define INSTRUCTION_PAGE_FAULT_TYPE       0xC
4  #define LOAD_PAGE_FAULT_TYPE              0xD
5  #define STORE_PAGE_FAULT_TYPE             0xF
6
7  void trap_handler(uint64_t scause, uint64_t sepc, struct pt_regs *regs) {
8      switch (scause) {
9          // Interrupts.
10         case SUPERVISOR_TIMER_INTERRUPT_TYPE:
11             ...
12         // Exceptions.
13         case ENVIRONMRNT_CALL_FROM_U_MODE_TYPE:
14             ...
15         case INSTRUCTION_PAGE_FAULT_TYPE:
16             [[fallthrough]];
17         case LOAD_PAGE_FAULT_TYPE:
18             [[fallthrough]];
19         case STORE_PAGE_FAULT_TYPE:
20             do_page_fault(regs);
21             break;
22         // default.
23         default:
24             ...
25     }
26 }

```

然后具体实现缺页异常处理，根据以下流程走

1. 检查是否是Bad Address，根据 `stval` 是否在进程的VMA所属范围中判断
2. 判断是否非法访问，根据VMA标注的权限和中断原因判断，其中
 1. `INSTRUCTION_PAGE_FAULT` 类型需要执行权限

2. `LOAD_PAGE_FAULT` 类型需要可读权限
3. `STORE_PAGE_FAULT` 类型需要可写权限
3. 合法访问，先分配一个页
4. 判断是否是匿名区域访问
 1. 如果是，直接映射到页表中
 2. 如果否，是用户程序数据，先拷贝后映射到页表

由于内存分布是从 `vaddr` 开始的，如果访问的地址处在 `vaddr` 所在的页内，就把这个页拷贝；如果是下一个页内的，就把下一个页拷贝

拷贝是从 `_sramdisk + pgoff` 开始

如果超过了 `filesz` 的部分，根据lab4，需要把超出的部分置0，因此如果是整页全超出了就不必进行拷贝了，这里需要叠加一个判断

C `arch/riscv/kernel/trap.c`

```

1  extern struct task_struct *current;
2  extern char _sramdisk[];
3  extern char _eramdisk[];
4  extern void create_mapping(uint64_t *pgtbl, uint64_t va, uint64_t pa, uint64_t
    sz, uint64_t perm);
5
6  void do_page_fault(struct pt_regs *regs) {
7      // Check bad address.
8      uint64_t addr = csr_read(stval);
9      struct vm_area_struct *vma = find_vma(&current->mm, addr);
10     if (!vma) error("Bad address for page fault at 0x%lx", addr);
11
12     // Check permission.
13     uint64_t perm = vma->vm_flags; // Get the permission of the vma.
14     uint64_t cause = csr_read(scause); // Get the cause of the trap.
15     bool is_illegal = (
16         cause == INSTRUCTION_PAGE_FAULT_TYPE && !(perm & VM_EXEC) ||
17         cause == LOAD_PAGE_FAULT_TYPE && !(perm & VM_READ) ||
18         cause == STORE_PAGE_FAULT_TYPE && !(perm & VM_WRITE)
19     );
20     if (is_illegal) error("Illegal access to 0x%lx", addr); // Permission
    denied.
21
22     // Legal access, continue.
23     uint64_t *page = (uint64_t *)alloc_page();
24     memset((void *)page, 0x0, PGSIZE);
25     if (!(perm & VM_ANON)) {
26         uint64_t *src = (uint64_t *)((uint64_t)_sramdisk + PGROUNDDOWN(vma-
    >vm_pgoff + addr - vma->vm_start));
27         if (addr ≥ PGROUNDDOWN(vma->vm_start + vma->vm_filesz) && addr <
    PGROUNDUP(vma->vm_start + vma->vm_filesz)) {
28             memcpy((void *)page, (void *)src, PGOFFSET(vma->vm_filesz));
29         } else if (addr < PGROUNDDOWN(vma->vm_start + vma->vm_filesz))
    memcpy((void *)page, (void *)src, PGSIZE);

```

```

30     }
31     create_mapping(current->pgd, PGROUNDNDOWN(addr), VA2PA((uint64_t)page),
    PGSIZE, 0b11010001 | (perm & (VM_READ | VM_WRITE | VM_EXEC)));
32 }

```

测试缺页处理

运行测试PFH1:

```

SET [PID = 3 PRIORITY = 4 COUNTER = 4]
SET [PID = 4 PRIORITY = 1 COUNTER = 1]
[PID = 0 PRIORITY = 0 COUNTER = 0] --> [PID = 2 PRIORITY = 10 COUNTER = 10]
[trap.c,38,do_page_fault] [PID = 2 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e1000, 0x802e2000) -> [0x10000, 0x11000), perm: df

[trap.c,38,do_page_fault] [PID = 2 PC = 0x10178] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e4000, 0x802e5000) -> [0x3ffffff00, 0x400000000), perm: d7

[trap.c,38,do_page_fault] [PID = 2 PC = 0x1019c] Valid page fault at 0x12000, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e7000, 0x802e8000) -> [0x12000, 0x13000), perm: df

[trap.c,38,do_page_fault] [PID = 2 PC = 0x110cc] Valid page fault at 0x110cc, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e8000, 0x802e9000) -> [0x11000, 0x12000), perm: df

[U-MODE] pid: 2, sp is 0x3fffffff0, this is print No.1
[U-MODE] pid: 2, sp is 0x3fffffff0, this is print No.2
[PID = 2 PRIORITY = 10 COUNTER = 0] --> [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,38,do_page_fault] [PID = 1 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802e9000, 0x802ea000) -> [0x10000, 0x11000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10178] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802ec000, 0x802ed000) -> [0x3ffffff00, 0x400000000), perm: d7

[trap.c,38,do_page_fault] [PID = 1 PC = 0x1019c] Valid page fault at 0x12000, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802ef000, 0x802f0000) -> [0x12000, 0x13000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x110cc] Valid page fault at 0x110cc, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802f0000, 0x802f1000) -> [0x11000, 0x12000), perm: df

[U-MODE] pid: 1, sp is 0x3fffffff0, this is print No.1

```

运行测试PFH2:

```

SET [PID = 3 PRIORITY = 4 COUNTER = 4]
SET [PID = 4 PRIORITY = 1 COUNTER = 1]
[PID = 0 PRIORITY = 0 COUNTER = 0] --> [PID = 2 PRIORITY = 10 COUNTER = 10]
[trap.c,38,do_page_fault] [PID = 2 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e1000, 0x802e2000] -> [0x10000, 0x11000), perm: df

[trap.c,38,do_page_fault] [PID = 2 PC = 0x10178] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e4000, 0x802e5000] -> [0x3fffffff000, 0x4000000000), perm: d7

[trap.c,38,do_page_fault] [PID = 2 PC = 0x10198] Valid page fault at 0x13000, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e7000, 0x802e8000] -> [0x13000, 0x14000), perm: df

[trap.c,38,do_page_fault] [PID = 2 PC = 0x110a4] Valid page fault at 0x110a4, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d4000, [0x802e8000, 0x802e9000] -> [0x11000, 0x12000), perm: df

[U-MODE] pid: 2, increment: 0
[U-MODE] pid: 2, increment: 1
[PID = 2 PRIORITY = 10 COUNTER = 0] --> [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,38,do_page_fault] [PID = 1 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802e9000, 0x802ea000] -> [0x10000, 0x11000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10178] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802ec000, 0x802ed000] -> [0x3fffffff000, 0x4000000000), perm: d7

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10198] Valid page fault at 0x13000, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802ef000, 0x802f0000] -> [0x13000, 0x14000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x110a4] Valid page fault at 0x110a4, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802f0000, 0x802f1000] -> [0x11000, 0x12000), perm: df

[U-MODE] pid: 1, increment: 0

```

实现 fork 系统调用

准备工作

设置全局变量 `nr_tasks`，仅初始化一个进程

添加系统调用号，扩展系统调用

拷贝内核栈

创建新进程的结构体，并直接复制当前进程的内核栈

C `arch/riscv/kernel/syscall.c`

```

1  extern uint64_t nr_tasks;
2  extern struct task_struct *task[NR_TASKS];
3  extern uint64_t swapper_pg_dir[];
4  extern void __ret_from_fork();
5  extern void create_mapping(uint64_t *pgtbl, uint64_t va, uint64_t pa, uint64_t
    sz, uint64_t perm);
6  uint64_t do_fork(struct pt_regs *regs) {
7      struct task_struct *_task = (struct task_struct *)alloc_page();
8      memcpy((void *)_task, (void *)current, PGSIZE);
9      ...
10 }

```

创建页表

首先先拷贝内核页表，直接拷贝即可

```
C arch/riscv/kernel/syscall.c/do_fork()

1  uint64_t do_fork(struct pt_regs *regs) {
2      ...
3      _task->pgd = (uint64_t *)alloc_page();
4      memset((void *)_task->pgd, 0x0, PGSIZE);
5      memcpy((void *)_task->pgd, (void *)swapper_pg_dir, PGSIZE);
6      ...
7  }
```

之后遍历父进程的VMA和页表

先实现一个在进程页表中进行查找的函数，用于查找目标地址是否在页表中存在，即是否PTE有效

查找的原理和创建三级页表原理一致，一级一级查找，如果这一级没找到就直接说明不存在，如果找到了就顺着往下一级查找

最终返回结果表示是否存在

```
C arch/riscv/kernel/syscall.c

1  bool find_pte(uint64_t *pgtbl, uint64_t va) {
2      uint64_t *pgtbl2 = pgtbl;
3      uint64_t vpn2 = (va >> 30) & 0x1FF;
4      if (!(pgtbl2[vpn2] & 0x1)) return false;
5
6      uint64_t *pgtbl1;
7      uint64_t vpn1 = (va >> 21) & 0x1FF;
8      pgtbl1 = (uint64_t *)((pgtbl2[vpn2] >> 10 << 12) + PA2VA_OFFSET);
9      if (!(pgtbl1[vpn1] & 0x1)) return false;
10
11     uint64_t *pgtbl0;
12     uint64_t vpn0 = (va >> 12) & 0x1FF;
13     pgtbl0 = (uint64_t *)((pgtbl1[vpn1] >> 10 << 12) + PA2VA_OFFSET);
14     return pgtbl0[vpn0] & 0x1;
15 }
```

完成之后进行父进程的VMA遍历

1. 获取父进程的mmap
2. 遍历mmap
 1. 每一个父进程中的VMA都使用 `do_mmap()` 完成在子进程中的添加
 2. 对于每一个VMA，从它的起始地址 `vm_start` 开始，一直按页递增到 `vm_end`，调用 `find_pte()` 查找是否在父进程页表中存在，如果存在，就拷贝整页内容，并映射到子进程页表中，注意这里的 `va` 已经是对齐了的，所以可以放心使用

```

C arch/riscv/kernel/syscall.c
1  uint64_t do_fork(struct pt_regs *regs) {
2      ...
3      _task->mm.mmap = NULL;
4      struct vm_area_struct *vma = current->mm.mmap;
5      while (vma) {
6          do_mmap(&_task->mm, vma->vm_start, vma->vm_end - vma->vm_start, vma->
>vm_pgoff, vma->vm_filesz, vma->vm_flags);
7          for (uint64_t va = PGROUNDDOWN(vma->vm_start); va ≤ vma->vm_end; va +=
PGSIZE) {
8              if(find_pte(current->pgd, va)) {
9                  uint64_t *page = (uint64_t *)alloc_page();
10                 memcpy((void *)page, (void *)va, PGSIZE);
11                 create_mapping(_task->pgd, va, VA2PA((uint64_t)page), PGSIZE,
0b11010001 | (vma->vm_flags & (VM_READ | VM_WRITE | VM_EXEC)));
12             }
13         }
14         vma = vma->vm_next;
15     }
16     ...
17 }

```

进程返回

父进程返回是直接通过 `do_fork()` 返回

对于子进程来说

- `ra` 设为 `__ret_from_fork`，子进程认为自己也是刚刚执行完一个中断处理，从 `__traps` 返回
- `sp` 设为当前子进程的内核栈指针
- `sscratch` 设为当前的 `sscratch` 寄存器值

对于子进程来说，它的 `struct pt_regs` 寄存器保存的地址与子进程的偏移量是和父进程与其 `struct pt_regs` 寄存器偏移量一致的，因为在拷贝内核栈时已经把寄存器的相关信息一起拷贝了，已经在子进程的内核栈中，所以直接计算出即可

子进程的 `sp` 设为刚刚的 `thread.sp`，并设返回值 `a0` 为0，并手动 `sepc+4`

最后添加到进程列表中，父进程返回PID

```

C arch/riscv/kernel/syscall.c
1  uint64_t do_fork(struct pt_regs *regs) {
2      ...
3      _task->thread.ra = (uint64_t)__ret_from_fork;
4      _task->thread.sp = (uint64_t)_task - (uint64_t)current + regs->
>regs[reg_sp];
5      _task->thread.sscratch = csr_read(sscratch);
6
7      struct pt_regs *child_regs = (struct pt_regs *)((uint64_t)_task -
(uint64_t)current + (uint64_t)regs);

```



```

8     child_regs->regs[reg_sp] = _task->thread.sp;
9     child_regs->regs[reg_a0] = 0;
10    child_regs->sepc += 4;
11
12    task[_task->pid = nr_tasks++] = _task; // Add to task table.
13    return _task->pid;
14 }

```

测试Fork

测试Fork1:

```

...setup_vm_final done.
...task_init done!
2024 ZJU Operating System
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
[PID = 0 PRIORITY = 0 COUNTER = 0] --> [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,38,do_page_fault] [PID = 1 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002ce000, [0x802d1000, 0x802d2000) -> [0x10000, 0x11000), perm: df
[trap.c,38,do_page_fault] [PID = 1 PC = 0x101ac] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002ce000, [0x802d4000, 0x802d5000) -> [0x3fffffff000, 0x4000000000), perm: d7
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d8000, [0x802da000, 0x802db000) -> [0x10000, 0x11000), perm: df
[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d8000, [0x802de000, 0x802df000) -> [0x3fffffff000, 0x4000000000), perm: d7
[syscall.c,77,do_fork] Fork page table done

[syscall.c,94,do_fork] [PID = 1] forked from [PID = 2]

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10240] Valid page fault at 0x12000, cause = 0xd, perm = 0xe
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002ce000, [0x802e1000, 0x802e2000) -> [0x12000, 0x13000), perm: df
[trap.c,38,do_page_fault] [PID = 1 PC = 0x1114c] Valid page fault at 0x1114c, cause = 0xc, perm = 0xe
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002ce000, [0x802e2000, 0x802e3000) -> [0x11000, 0x12000), perm: df

[U-PARENT] pid: 1 is running! global_variable: 0
[PID = 1 PRIORITY = 7 COUNTER = 0] --> [PID = 2 PRIORITY = 7 COUNTER = 7]
[trap.c,38,do_page_fault] [PID = 2 PC = 0x101e8] Valid page fault at 0x12000, cause = 0xd, perm = 0xe
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d8000, [0x802e3000, 0x802e4000) -> [0x12000, 0x13000), perm: df
[trap.c,38,do_page_fault] [PID = 2 PC = 0x1114c] Valid page fault at 0x1114c, cause = 0xc, perm = 0xe
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002d8000, [0x802e4000, 0x802e5000) -> [0x11000, 0x12000), perm: df

[U-CHILD] pid: 2 is running! global_variable: 0
[U-CHILD] pid: 2 is running! global_variable: 1
SET [PID = 1 PRIORITY = 7 COUNTER = 7]

```

测试Fork2:

```

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10470] Valid page fault at 0x14008, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802da000, 0x802db000) -> [0x14000, 0x15000), perm: df

[U] pid: 1 is running! global_variable: 0
[U] pid: 1 is running! global_variable: 1
[U] pid: 1 is running! global_variable: 2
[trap.c,38,do_page_fault] [PID = 1 PC = 0x10230] Valid page fault at 0x13008, cause = 0xf, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802db000, 0x802dc000) -> [0x13000, 0x14000), perm: df

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802df000, 0x802e0000) -> [0x10000, 0x11000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802e2000, 0x802e3000) -> [0x11000, 0x12000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802e3000, 0x802e4000) -> [0x12000, 0x13000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802e4000, 0x802e5000) -> [0x13000, 0x14000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802e5000, 0x802e6000) -> [0x14000, 0x15000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802e7000, 0x802e8000) -> [0x3fffffff000, 0x40000000000), perm: d7

[syscall.c,77,do_fork] Fork page table done

[syscall.c,94,do_fork] [PID = 1] forked from [PID = 2]

[U-PARENT] pid: 1 is running! Message: ZJU OS Lab5
[U-PARENT] pid: 1 is running! global_variable: 3
[PID = 1 PRIORITY = 7 COUNTER = 0] --> [PID = 2 PRIORITY = 7 COUNTER = 7]
[U-CHILD] pid: 2 is running! Message: ZJU OS Lab5
[U-CHILD] pid: 2 is running! global_variable: 3
[U-CHILD] pid: 2 is running! global_variable: 4

```

测试Fork3:

```
[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff000316000, [0x8031e000, 0x8031f000) -> [0xffffffff000, 0x4000000000), perm: d7

[syscall.c,77,do_fork] Fork page table done

[syscall.c,94,do_fork] [PID = 3] forked from [PID = 7]

[U] pid: 3 is running! global_variable: 2
[U] pid: 3 is running! global_variable: 3
[PID = 3 PRIORITY = 7 COUNTER = 0] --> [PID = 5 PRIORITY = 7 COUNTER = 7]
[U] pid: 5 is running! global_variable: 1
[vm.c,53,create_mapping] pgtbl: 0xffffffff000322000, [0x80324000, 0x80325000) -> [0x10000, 0x11000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff000322000, [0x80327000, 0x80328000) -> [0x11000, 0x12000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff000322000, [0x80328000, 0x80329000) -> [0x12000, 0x13000), perm: df

[syscall.c,77,do_fork] Fork page table done

[vm.c,53,create_mapping] pgtbl: 0xffffffff000322000, [0x8032a000, 0x8032b000) -> [0xffffffff000, 0x4000000000), perm: d7

[syscall.c,77,do_fork] Fork page table done

[syscall.c,94,do_fork] [PID = 5] forked from [PID = 8]

[U] pid: 5 is running! global_variable: 2
[U] pid: 5 is running! global_variable: 3
[PID = 5 PRIORITY = 7 COUNTER = 0] --> [PID = 6 PRIORITY = 7 COUNTER = 7]
[U] pid: 6 is running! global_variable: 2
[U] pid: 6 is running! global_variable: 3
[PID = 6 PRIORITY = 7 COUNTER = 0] --> [PID = 7 PRIORITY = 7 COUNTER = 7]
[U] pid: 7 is running! global_variable: 2
[U] pid: 7 is running! global_variable: 3
[PID = 7 PRIORITY = 7 COUNTER = 0] --> [PID = 8 PRIORITY = 7 COUNTER = 7]
[U] pid: 8 is running! global_variable: 2
[U] pid: 8 is running! global_variable: 3
[PID = 8 PRIORITY = 7 COUNTER = 0] --> [PID = 4 PRIORITY = 7 COUNTER = 6]
[U] pid: 4 is running! global_variable: 2
```

Exercises

1 呈现出你在 page fault 的时候拷贝 ELF 程序内容的逻辑

拷贝ELF程序时

1. 分配一个页 `page` 用于拷贝内容
2. 判断不是匿名区域
3. 计算出拷贝的源地址 `uint64_t *src = (uint64_t *)((uint64_t)_sramdisk + PGROUNDDOWN(vma->vm_pgoff + addr - vma->vm_start))`
4. 判断Bad Address是否超出 `vm_start + vm_filesz`
 1. 如果 `addr` 所在页地址小于 `vm_start + vm_filesz` 所在页，则直接拷贝即可
 2. 如果 `addr` 所在页刚好就是 `vm_start + vm_filesz` 所在页，则仅拷贝小于的那部分，超出的部分因为在分配页时已经初始化置0所以不用额外操作
 3. 如果 `addr` 所在页地址大于 `vm_start + vm_filesz` 所在页，则无需拷贝
5. 把 `addr` 所在页低地址映射到页表中，设置权限值

其中，拷贝的源地址是从 `(uint64_t)_sramdisk + vma->vm_pgoff` 开始，目标地址是 `addr`，而且这里是按页拷贝，也就是说要得到 `addr` 所在的那一页

🔗 切入点

如果 `addr == vma->vm_start` , 那么就是拷贝第一页, 即从 `(uint64_t)_sramdisk` 拷贝 `PGSIZE` 内容到 `PGROUNDOWN(addr)` 上

Plain text 拷贝ELF示意图



C arch/riscv/kernel/trap.c

```

1 void do_page_fault(struct pt_regs *regs) {
2     // Check bad address.
3     uint64_t addr = csr_read(stval);
4     struct vm_area_struct *vma = find_vma(&current->mm, addr);
5     if (!vma) error("Bad address for page fault at 0x%lx", addr);
6
7     // Check permission.
8     uint64_t perm = vma->vm_flags; // Get the permission of the vma.
9     uint64_t cause = csr_read(scause); // Get the cause of the trap.
10    bool is_illegal = (
11        cause == INSTRUCTION_PAGE_FAULT_TYPE && !(perm & VM_EXEC) ||
12        cause == LOAD_PAGE_FAULT_TYPE && !(perm & VM_READ) ||
13        cause == STORE_PAGE_FAULT_TYPE && !(perm & VM_WRITE)
14    );
15    if (is_illegal) error("Illegal access to 0x%lx", addr); // Permission
    denied.
16
17    // Legal access, continue.
18    uint64_t *page = (uint64_t *)alloc_page();

```

```

19     memset((void *)page, 0x0, PGSIZE);
20     if (!(perm & VM_ANON)) {
21         uint64_t *src = (uint64_t *)((uint64_t)_sramdisk + PGROUNDDOWN(vma->vm_pgoff + addr - vma->vm_start));
22         if (addr ≥ PGROUNDDOWN(vma->vm_start + vma->vm_filesz) && addr < PGROUNDUP(vma->vm_start + vma->vm_filesz)) {
23             memcpy((void *)page, (void *)src, PGOFFSET(vma->vm_filesz));
24         } else if (addr < PGROUNDDOWN(vma->vm_start + vma->vm_filesz))
25             memcpy((void *)page, (void *)src, PGSIZE);
26         create_mapping(current->pgd, PGROUNDDOWN(addr), VA2PA((uint64_t)page), PGSIZE, 0b11010001 | (perm & (VM_READ | VM_WRITE | VM_EXEC)));
27     }

```

2 回答 4.3.5 中的问题

2.1 在 do_fork 中，父进程的内核栈和用户栈指针分别是什么

在中断处理 `__traps` 中，先是把用户态栈和内核态栈进行了切换，原来的用户态栈指针 `sp` 保存在了 `sscratch` 中，而内核态栈的指针从 `sscratch` 中读出，存入 `sp` 中，而这个 `sp` 又保存在了 `struct pt_regs` 这个结构体中用于给 `trap_handler` 传参，一路由 `trap_handler` 传到 `syscall` 传到 `do_fork` 中，所以

- 内核栈指针: `regs[reg_sp]`
- 用户栈指针: `sscratch`

2.2 在 do_fork 中，子进程的内核栈和用户栈指针的值应该是什么

- 内核栈指针值应该等于内核栈上保存的寄存器 `sp` 的值
- 用户栈指针值应该等于 `sscratch` 的值

2.3 在 do_fork 中，子进程的内核栈和用户栈指针分别应该赋值给谁

- 内核栈指针应该赋值给内核栈上的 `struct pt_regs` 中的 `sp` 寄存器
- 用户栈指针应该赋值给子进程的 `thread.sscratch`

这样，在子进程经过 `schedule()` 调度，经过 `__switch_to` 切换后从 `__traps` 中返回时有一个恢复上下文的过程，在这里会有二者的交换，交换后就是正确的

3 为什么要为子进程 pt_regs 的 sepc 手动加四？

因为子进程返回并没有从 `syscall()` 里面返回，因此需要在 `Fork` 时手动加四，使得中断返回正常执行用户态程序后执行的下一条指令是 `Fork` 后的指令

4 对于 Fork main#2 (即 FORK2)，在运行时，ZJU OS Lab5 位于内存的什么位置？是否在读取的时候产生了 page fault？请给出必要的截图以说明

如下图所述，运行时在内存的 `0x13008` 所在页，产生了缺页异常

```

...setup_vm_final done.
...task_init done!
2024 ZJU Operating System
SET [PID = 1 PRIORITY = 7 COUNTER = 7]
[PID = 0 PRIORITY = 0 COUNTER = 0] --> [PID = 1 PRIORITY = 7 COUNTER = 7]
[trap.c,38,do_page_fault] [PID = 1 PC = 0x100e8] Valid page fault at 0x100e8, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802d2000, 0x802d3000] -> [0x10000, 0x11000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x101ac] Valid page fault at 0x3fffffff8, cause = 0xf, perm = 0x7

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802d5000, 0x802d6000] -> [0x3fffffff000, 0x40000000000), perm: d7

[trap.c,38,do_page_fault] [PID = 1 PC = 0x101d4] Valid page fault at 0x12000, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802d8000, 0x802d9000] -> [0x12000, 0x13000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x11320] Valid page fault at 0x11320, cause = 0xc, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802d9000, 0x802da000] -> [0x11000, 0x12000), perm: df

[trap.c,38,do_page_fault] [PID = 1 PC = 0x10470] Valid page fault at 0x14008, cause = 0xd, perm = 0xe

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802da000, 0x802db000] -> [0x14000, 0x15000), perm: df

[U] pid: 1 is running! global_variable: 0
[U] pid: 1 is running! global_variable: 1
[U] pid: 1 is running! global_variable: 2
[trap.c,38,do_page_fault] [PID = 1 PC = 0x10230] Valid page fault at 0x13008, cause = 0xf, perm = 0xe
[vm.c,53,create_mapping] pgtbl: 0xffffffff0002cf000, [0x802db000, 0x802dc000] -> [0x13000, 0x14000), perm: df

[syscall.c,57,do_fork] Forking...

[syscall.c,77,do_fork] Fork page table start ↓↓↓

[vm.c,53,create_mapping] pgtbl: 0xffffffff0002dd000, [0x802df000, 0x802e0000] -> [0x10000, 0x11000), perm: df

[syscall.c,79,do_fork] Fork page table done ↑↑↑

[syscall.c,77,do_fork] Fork page table start ↓↓↓

```

5 画图分析 make run TEST=FORK3 的进程 fork 过程，并呈现出各个进程的 `global_variable` 应该从几开始输出，再与你的输出进行对比验证

Fork3的用户程序如下，为三个Fork添加了标号

C `user/main.c`

```

1  int global_variable = 0;
2
3  int main() {
4
5      printf("[U] pid: %ld is running! global_variable: %d\n", getpid(),
        global_variable++);
6      fork(); // Fork 1.
7      fork(); // Fork 2.
8
9      printf("[U] pid: %ld is running! global_variable: %d\n", getpid(),
        global_variable++);
10     fork(); // Fork 3.
11
12     while(1) {
13         printf("[U] pid: %ld is running! global_variable: %d\n", getpid(),
            global_variable++);
14         wait(WAIT_TIME);
15     }

```

初始时只有PID=1被初始化并运行

Plain text

```

1
2  PID          1  4  3  7  2  6  5  8
3
4  [PID 1] ==>  T
5  |
6  print(0)    |
7  |
8  Fork 1  ->  |-----|
9  |           |
10 |           |
11 |           |
12 |           |
13 |           |
14 |           |
15 |           |
16 [PID 2] ==>  T
17 |           |
18 |           |
19 |           |
20 |           |
21 |           |
22 |           |
23 |           |
24 [PID 3] ==>  T
25 |           |
26 |           |
27 |           |
28 |           |
29 |           |
30 [PID 5] ==>  T
31 |           |
32 |           |
33 |           |
34 |           |
35 |           |
36 [PID 6] ==>  T
37 |           |
38 |           |
39 |           |
40 [PID 7] ==>  T
41 |           |
42 |           |
43 |           |
44 [PID 8] ==>  T
45 |           |
46 |           |

```

```

47      *   *   *   *   *   *   *   |
48  [PID 4] ==> *   T   *   *   *   *   *   ⊥
49      *   |   *   *   *   *   *   *
50  print(2)    *   |   *   *   *   *   *   *
51      *   |   *   *   *   *   *   *
52  [PID 1] ==> T   ⊥   *   *   *   *   *   *
53      |   *   *   *   *   *   *   *
54
55      ...
56

```

Thinkings

这个实验还是蛮有内容的，花了很多时间，因为debug比较麻烦所以耗时好多天才完成，最后还花了一些时间做实验报告，特别是两个画图都比较花时间，但是总体上对缺页异常和Fork系统调用理解更深了